


PRACTICAL NO : 1

Aim: Install, configure and run Hadoop and HDFS

Hadoop Installation.

Step 1: download java jdk first .the package size 168.67MB

Windows x64	168.67 MB	 jdk-8u291-windows-x64.exe
-------------	-----------	---

hadoop-2.10.1-src.tar.gz	16-05-2021 17:16	WinRAR archive	43,967 KB
hqbhib.txt	06-05-2021 08:23	Text Document	1 KB
jdk-8u291-windows-x64.exe	16-05-2021 17:16	Application	1,72,731 KB
LogisticRegressionGFG.png	23-05-2021 17:04	PNG File	4 KB

Step 2: download Hadoop binaries from the official website. The binary package size is about 342 MB.

Download

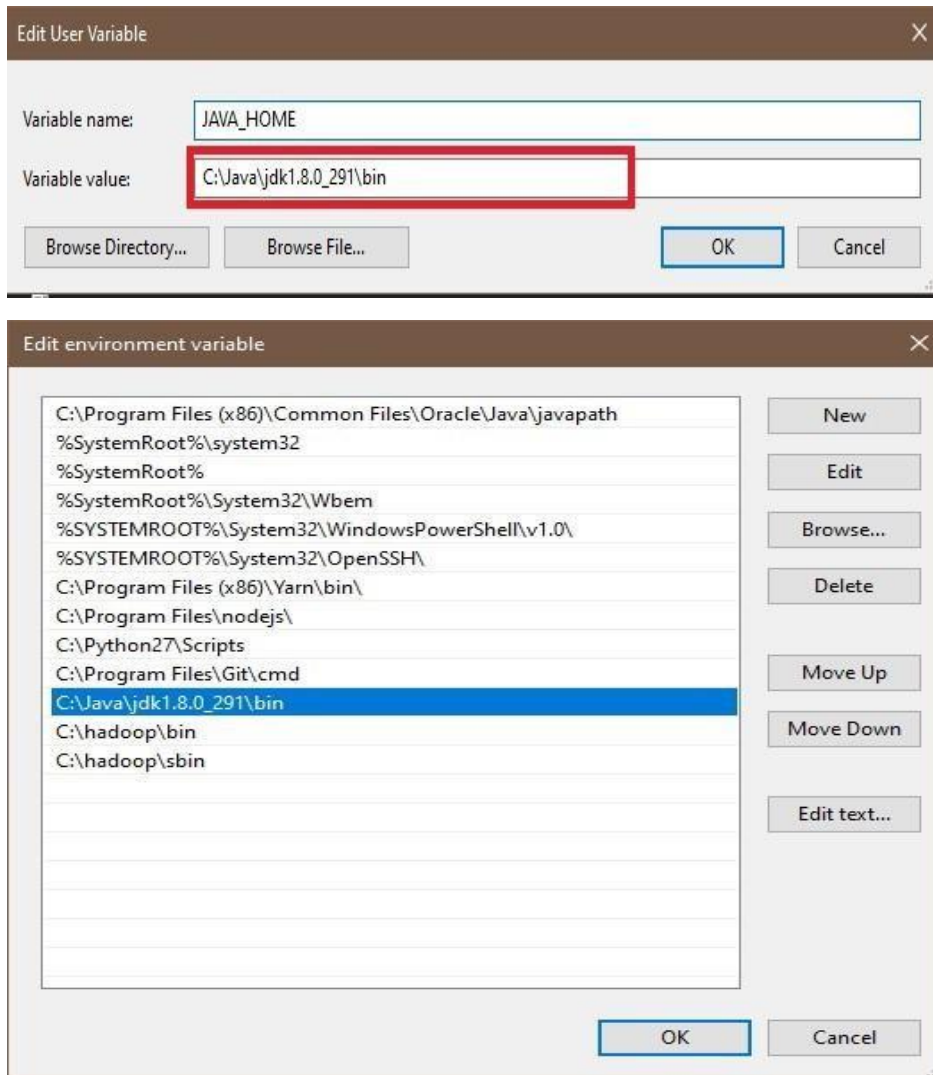
Hadoop is released as source code tarballs with corresponding binary tarballs for convenience. The downloads are distributed via mirror sites and should be checked for tampering using GPG or SHA-512.

Version	Release date	Source download	Binary download	Release notes
3.2.2	2021 Jan 9	source (checksum signature)	binary (checksum signature)	Announcement
2.10.1	2020 Sep 21	source (checksum signature)	binary (checksum signature)	Announcement
3.1.4	2020 Aug 3	source (checksum signature)	binary (checksum signature)	Announcement
3.3.0	2020 Jul 14	source (checksum signature)	binary (checksum signature) binary-aarch64 (checksum signature)	Announcement

Step 3: After finishing the file download, we should unpack the package using 7zip in two steps. First, we should extract the hadoop -3.2.1 .tar.gz library, and then, we should unpack the extracted tar file:

Name	Date modified	Type	Size
hadoop-3.3.0.tar.gz	12-05-2021 08:51	WinRAR archive	4,89,013 KB
wavelets_0.3-0.2.tar.gz	12-05-2021 08:27	WinRAR archive	114 KB
govind.data	12-05-2021 08:24	DATA File	283 KB

Step 4: When the “Advanced system settings” dialog appears, go to the “Advanced” tab and click on the “Environment variables” button located on the bottom of the dialog.



Step 5: Check the version of java

```

C:\Windows\system32\cmd.exe
Microsoft Windows [Version 10.0.19041.928]
(c) Microsoft Corporation. All rights reserved.

C:\Users\hp>javac
Usage: javac <options> <source files>
where possible options include:
  -g               Generate all debugging info
  -g:none          Generate no debugging info
  -g:{lines,vars,source}  Generate only some debugging info
  -nowarn          Generate no warnings
  -verbose         Output messages about what the compiler is doing
  -deprecation     Output source locations where deprecated APIs are used
  -classpath <path> Specify where to find user class files and annotation processors
  -cp <path>       Specify where to find user class files and annotation processors
  -sourcepath <path> Specify where to find input source files
  -bootclasspath <path> Override location of bootstrap class files
  -extdirs <dirs>   Override location of installed extensions
  -endorseddirs <dirs> Override location of endorsed standards path
  -proc:{none,only} Control whether annotation processing and/or compilation is done
  -processor <class1>[,<class2>,<class3>...] Names of the annotation processors to run; by

```

Step 6: Configuration core-site.xml

container-executor.cfg	07-07-2020 01:03	CFG File
core-site.xml	19-05-2021 17:57	XML File
hadoop-env.cmd	19-05-2021 17:57	Windows Command File

```

core-site.xml
C:\> hadoop > etc > hadoop > core-site.xml
1  <?xml version="1.0" encoding="UTF-8"?>
2  <?xml-stylesheet type="text/xsl" href="configuration.xsl"?>
3
4  <configuration>
5
6  <property>
7    <name>fs.defaultFS</name>
8    <value>hdfs://localhost:9000</value>
9  </property>
10 </configuration>

```

Step 7: Configuration core-site.xml

hdfs-rbf-site.xml	07-07-2020 00:26	XML File
hdfs-site.xml	19-05-2021 17:58	XML File
https-env.sh	07-07-2020 00:25	Shell Script

```

core-site.xml • hdfs-site.xml •
C: > hadoop > etc > hadoop > hdfs-site.xml
1  <?xml version="1.0" encoding="UTF-8"?>
2  <?xml-stylesheet type="text/xsl" href="configuration.xsl"?>
3
4  <configuration>
5  <property>
6      <name>dfs.replication</name>
7      <value>1</value>
8  </property>
9  <property>
10     <name>dfs.namenode.name.dir</name>
11     <value>C:\hadoop\data\namenode</value>
12
13 </property>
14 <property>
15     <name>dfs.namenode.data.dir</name>
16     <value>C:\hadoop\data\datanode</value>
17 </property>
18 </configuration>

```

Step 8: Configuration core-site.xml

mapred-queues.xml.template	07-07-2020 01:04	TEMPLATE File
mapred-site.xml	19-05-2021 17:58	XML File
ssl-client.xml.example	07-07-2020 00:16	EXAMPLE File

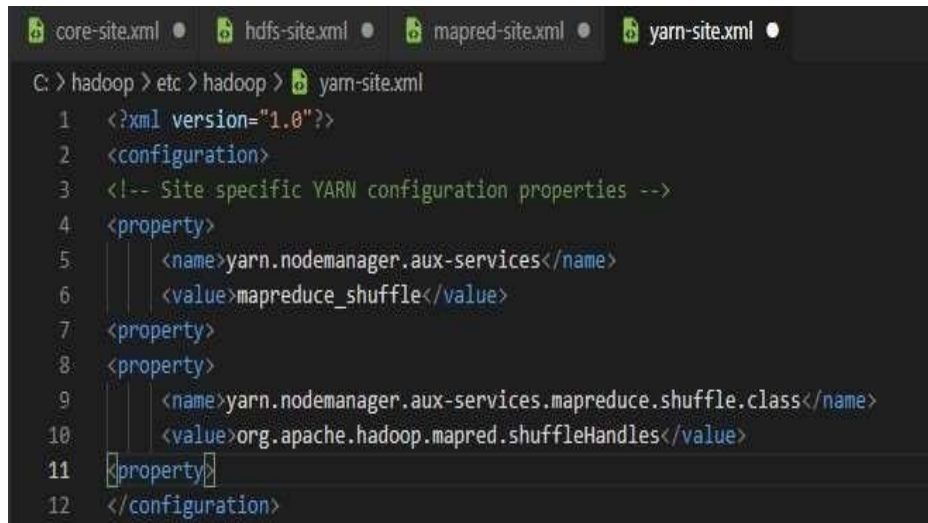
```

File Edit Selection View Go Run Terminal Help
core-site.xml • hdfs-site.xml • mapred-site.xml •
C: > hadoop > etc > hadoop > mapred-site.xml
1  <?xml version="1.0"?>
2  <?xml-stylesheet type="text/xsl" href="configuration.xsl"?>
3
4  <configuration>
5  <property>
6      <name>mapreduce.framework.name</name>
7      <value>yarn</value>
8  </configuration>

```

Step 9: Configuration core-site.xml

yarnservice-log4j.properties	07-07-2020 01:03	PROPERTIES File
yarn-site.xml	19-05-2021 17:58	XML File

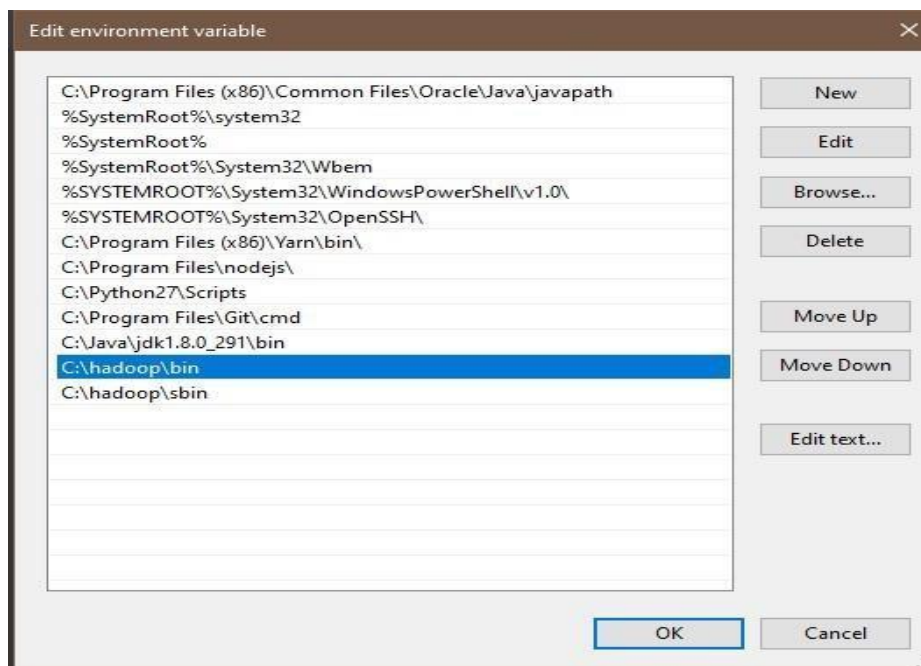


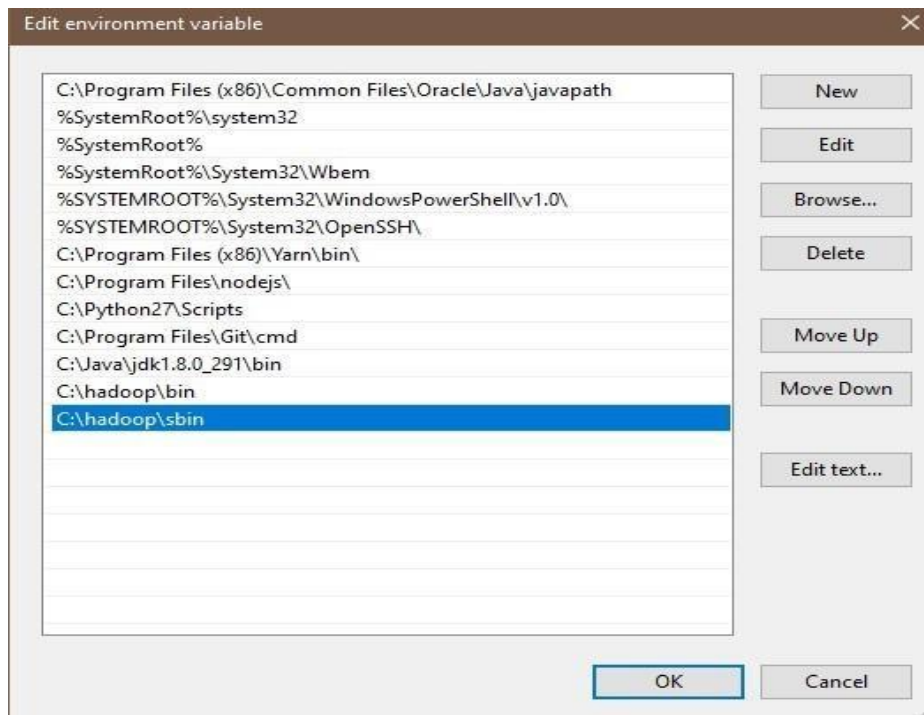
```

C:\> hadoop > etc > hadoop > yarn-site.xml
1  <?xml version="1.0"?>
2  <configuration>
3  <!-- Site specific YARN configuration properties -->
4  <property>
5      <name>yarn.nodemanager.aux-services</name>
6      <value>mapreduce_shuffle</value>
7  </property>
8  <property>
9      <name>yarn.nodemanager.aux-services.mapreduce.shuffle.class</name>
10     <value>org.apache.hadoop.mapred.shuffleHandles</value>
11 </property>
12 </configuration>

```

Step 10: When the “Advanced system settings” dialog appears, go to the “Advanced” tab and click on the “Environment variables” button located on the bottom of the dialog.





Step 11: let's check Hadoop install Successfully

```

C:\Windows\system32\cmd.exe
Java(TM) SE Runtime Environment (build 1.8.0_291-b10)
Java HotSpot(TM) 64-Bit Server VM (build 25.291-b10, mixed mode)

C:\Users\hp>hdfs namenode -format
2021-05-23 17:17:11,111 INFO namenode.NameNode: STARTUP_MSG:
/*****
STARTUP_MSG: Starting NameNode
STARTUP_MSG:   host = DESKTOP-VUUFK2Q/192.168.0.104
STARTUP_MSG:   args = [-format]
STARTUP_MSG:   version = 3.3.0
STARTUP_MSG:   classpath = C:\hadoop\etc\hadoop;C:\hadoop\share\hadoop\common;C:\h
s-smart-1.2.jar;C:\hadoop\share\hadoop\common\lib\animal-sniffer-annotations-1.17.
asm-5.0.4.jar;C:\hadoop\share\hadoop\common\lib\audience-annotations-0.5.0.jar;C:\
7.7.jar;C:\hadoop\share\hadoop\common\lib\checker-qual-2.5.2.jar;C:\hadoop\share\h
.4.jar;C:\hadoop\share\hadoop\common\lib\commons-cli-1.2.jar;C:\hadoop\share\hadoo
\hadoop\share\hadoop\common\lib\commons-collections-3.2.2.jar;C:\hadoop\share\hadoo
r;C:\hadoop\share\hadoop\common\lib\commons-configuration2-2.1.1.jar;C:\hadoop\sha
0.13.jar;C:\hadoop\share\hadoop\common\lib\commons-io-2.5.jar;C:\hadoop\share\hadoo
\hadoop\share\hadoop\common\lib\commons-logging-1.1.3.jar;C:\hadoop\share\hadoop\c
adoop\share\hadoop\common\lib\commons-net-3.6.jar;C:\hadoop\share\hadoop\common\li
\hadoop\common\lib\curator-client-4.2.0.jar;C:\hadoop\share\hadoop\common\lib\cura
e\hadoop\common\lib\curator-recipes-4.2.0.jar;C:\hadoop\share\hadoop\common\lib\dr
\common\lib\failureaccess-1.0.jar;C:\hadoop\share\hadoop\common\lib\gson-2.2.4.jar
va-27.0-jre.jar;C:\hadoop\share\hadoop\common\lib\hadoop-annotations-3.3.0.jar;C:\
auth-3.3.0.jar;C:\hadoop\share\hadoop\common\lib\hadoop-shaded-protobuf_3_7-1.0.0.
htrace-core4-4.1.0-incubating.jar;C:\hadoop\share\hadoop\common\lib\httpclient-4.5
ib\httpcore-4.4.10.jar;C:\hadoop\share\hadoop\common\lib\j2objc-annotations-1.1.ja

```

```

C:\Apache Hadoop Distribution
DEPRECATED: Use of this script to execute hdfs command is deprecated.
Instead use the hdfs command for it.
2021-05-23 17:19:33,116 INFO namenode.NameNode: STARTUP_MSG:
/*****
STARTUP_MSG: Starting NameNode
STARTUP_MSG: host = DESKTOP-VUUFK2Q/192.168.0.104
STARTUP_MSG: args = []
STARTUP_MSG: version = 3.3.0
STARTUP_MSG: classpath = C:\hadoop\etc\hadoop;C:\hadoop\share\hadoop\common;C:\hadoop\share\hadoop\common\lib\accessor
s-smart-1.2.jar;C:\hadoop\share\hadoop\common\lib\animal-sniffer-annotations-1.17.jar;C:\hadoop\share\hadoop\common\lib\
asm-5.0.4.jar;C:\hadoop\share\hadoop\common\lib\audience-annotations-0.5.0.jar;C:\hadoop\share\hadoop\common\lib\avro-1.
7.7.jar;C:\hadoop\share\hadoop\common\lib\checker-qual-2.5.2.jar;C:\hadoop\share\hadoop\common\lib\commons-beanutils-1.9
.4.jar;C:\hadoop\share\hadoop\common\lib\commons-cli-1.2.jar;C:\hadoop\share\hadoop\common\lib\commons-codec-1.11.jar;C:
\hadoop\share\hadoop\common\lib\commons-collections-3.2.2.jar;C:\hadoop\share\hadoop\common\lib\commons-compress-1.19.ja
r;C:\hadoop\share\hadoop\common\lib\commons-configuration2-2.1.1.jar;C:\hadoop\share\hadoop\common\lib\commons-daemon-1.
0.13.jar;C:\hadoop\share\hadoop\common\lib\commons-io-2.5.jar;C:\hadoop\share\hadoop\common\lib\commons-lang3-3.7.jar;C:
\hadoop\share\hadoop\common\lib\commons-logging-1.1.3.jar;C:\hadoop\share\hadoop\common\lib\commons-math3-3.1.1.jar;C:\h
adoop\share\hadoop\common\lib\commons-net-3.6.jar;C:\hadoop\share\hadoop\common\lib\commons-text-1.4.jar;C:\hadoop\share
\hadoop\common\lib\curator-client-4.2.0.jar;C:\hadoop\share\hadoop\common\lib\curator-framework-4.2.0.jar;C:\hadoop\shar
e\hadoop\common\lib\curator-recipes-4.2.0.jar;C:\hadoop\share\hadoop\common\lib\dnsjava-2.1.7.jar;C:\hadoop\share\hadoop

```

```

C:\Apache Hadoop Distribution
at com.ctc.wstx.sr.StreamScanner.throwParseError(StreamScanner.java:491)
at com.ctc.wstx.sr.StreamScanner.throwParseError(StreamScanner.java:475)
at com.ctc.wstx.sr.BasicStreamReader.reportWrongEndElem(BasicStreamReader.java:3365)
at com.ctc.wstx.sr.BasicStreamReader.readEndElem(BasicStreamReader.java:3292)
at com.ctc.wstx.sr.BasicStreamReader.nextFromTree(BasicStreamReader.java:2911)
at com.ctc.wstx.sr.BasicStreamReader.next(BasicStreamReader.java:1123)
at org.apache.hadoop.conf.Configuration$Parser.parseNext(Configuration.java:3347)
at org.apache.hadoop.conf.Configuration$Parser.parse(Configuration.java:3141)
at org.apache.hadoop.conf.Configuration.loadResource(Configuration.java:3034)
... 9 more

```

Step 12: Let check bin

```

C:\Windows\system32\cmd.exe

C:\Users\hp>cd C:\hadoop\sbin

C:\hadoop\sbin>start-all.cmd
This script is Deprecated. Instead use start-dfs.cmd and start-yarn.cmd
starting yarn daemons

C:\hadoop\sbin>

```

PRACTICAL NO : 2

Aim: Implement Decision tree classification techniques.

Step 1: The package "party" has the function ctree() which is used to create and analyze decision tree.

```
> install.packages("party")
```

Step 2: Load the party package. It will automatically load other# dependent packages. Print some records from data set readingsSkills.

```
> library("party")
> print(head(readingsSkills))
  nativespeaker age shoeSize    score
1          yes   5  24.83189 32.29385
2          yes   6  25.95238 36.63105
3           no  11  30.42170 49.60593
4          yes   7  28.66450 40.28456
5          yes  11  31.88207 55.46085
6          yes  10  30.07843 52.83124
>
```

Step 3 : Call function ctree to build a decision tree. The first parameter is a formula, which defines a target variable and a list of independent variables.

```
> library("party")
> str(iris)
'data.frame': 150 obs. of 5 variables:
 $ Sepal.Length: num  5.1 4.9 4.7 4.6 5 5.4 4.6 5 4.4 4.9 ...
 $ Sepal.Width : num  3.5 3 3.2 3.1 3.6 3.9 3.4 3.4 2.9 3.1 ...
 $ Petal.Length: num  1.4 1.4 1.3 1.5 1.4 1.7 1.4 1.5 1.4 1.5 ...
 $ Petal.Width : num  0.2 0.2 0.2 0.2 0.2 0.4 0.3 0.2 0.2 0.1 ...
 $ Species : Factor w/ 3 levels "setosa","versicolor",...: 1 1 1 1 1 1 1 1 1 1
```



```
> iris_ctree <- ctree(Species ~ Sepal.Length + Sepal.Width + Petal.Length + Petal.Width, data=iris)
> print(iris_ctree)
```

Conditional inference tree with 4 terminal nodes


Response: Species

Inputs: Sepal.Length, Sepal.Width, Petal.Length, Petal.Width

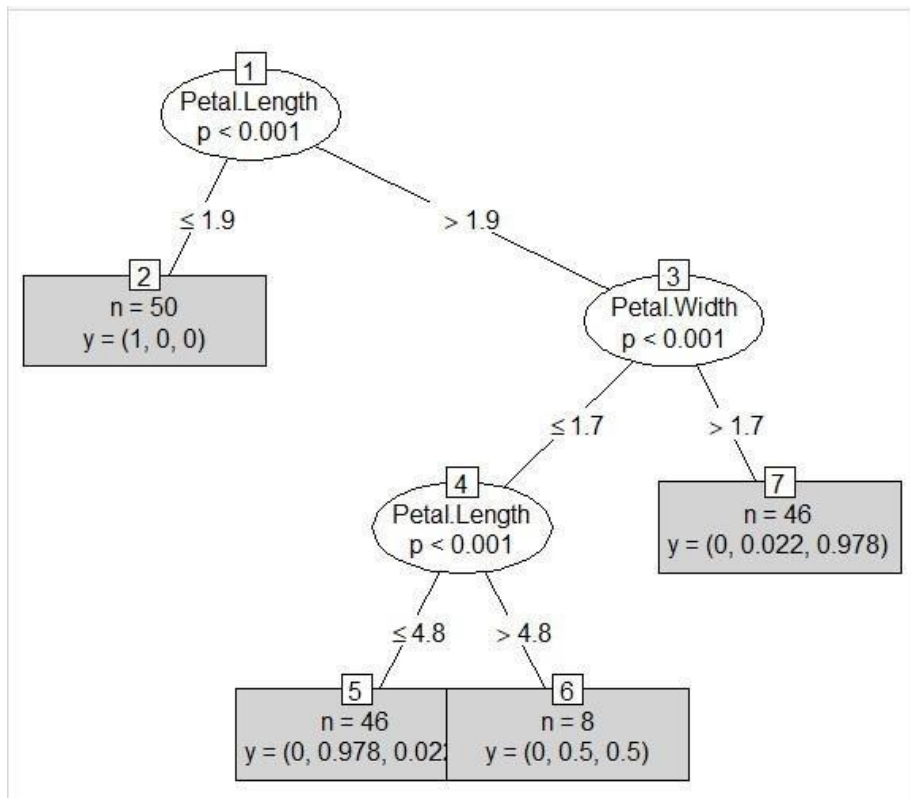
Number of observations: 150

```
1) Petal.Length <= 1.9; criterion = 1, statistic = 140.264
  2)* weights = 50
  1) Petal.Length > 1.9
    3) Petal.Width <= 1.7; criterion = 1, statistic = 67.894
      4) Petal.Length <= 4.8; criterion = 0.999, statistic = 13.865
        5)* weights = 46
        4) Petal.Length > 4.8
          6)* weights = 8
          3) Petal.Width > 1.7
            7)* weights = 46
> plot(iris_ctree)
```

Output :



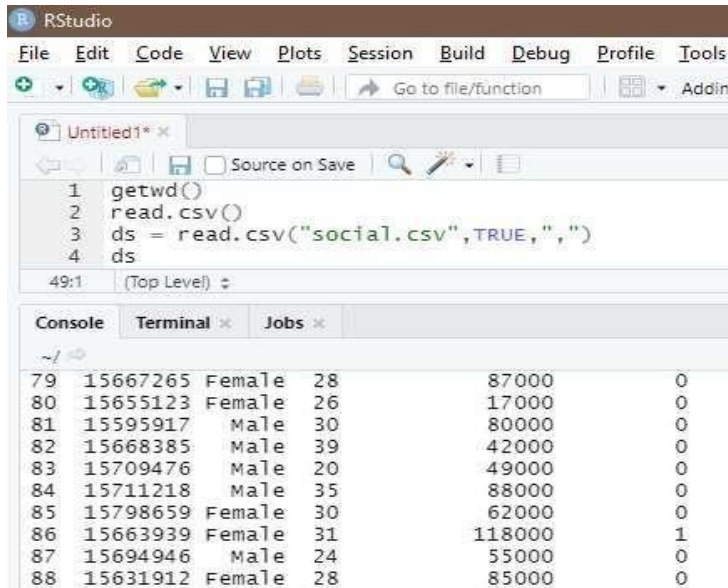
```
> plot(iris_ctree, type="simple")
```



PRACTICAL NO : 3

Aim: Classification using SVM

Step 1: Importing the dataset



```

1 getwd()
2 read.csv()
3 ds = read.csv("social.csv", TRUE, ",")
4 ds

```

49:1 (Top Level) ↕

79	15667265	Female	28	87000	0
80	15655123	Female	26	17000	0
81	15595917	Male	30	80000	0
82	15668385	Male	39	42000	0
83	15709476	Male	20	49000	0
84	15711218	Male	35	88000	0
85	15798659	Female	30	62000	0
86	15663939	Female	31	118000	1
87	15694946	Male	24	55000	0
88	15631912	Female	28	85000	0

Step 2: Selecting columns 3-5

```

> ds = ds[3:5]
> ds[3:5]
Error in `[.data.frame`(ds, 3:5) : undefined columns selected
> ds

```

	Age	Estimatedsalary	Purchased
1	19	19000	0
2	35	20000	0
3	26	43000	0
4	27	57000	0
5	19	76000	0
6	27	58000	0
7	27	84000	0
8	32	150000	1
9	25	33000	0
10	35	65000	0
11	26	80000	0
12	26	52000	0

Step 3: install package

```
> install.packages("caTools")
```

Step 4: Splitting the dataset

```
> library(caTools)
> set.seed(123)
> split = sample.split(ds$Purchased, splitRatio = 0.75)
> training_set = subset(ds, split == TRUE)
> test_set = subset(ds, split == FALSE)
> ds
```

	Age	EstimatedSalary	Purchased
1	19	19000	0
2	35	20000	0
3	26	43000	0
4	27	57000	0
5	19	76000	0
6	27	58000	0
7	27	84000	0
8	32	150000	1
9	25	33000	0
10	35	65000	0

Step 5: Feature Scaling

```
332 48      119000      1
333 42      65000      0
[ reached 'max' / getOption("max.print") -- omitted 67 rows ]
> test_set[-3] = scale(test_set[-3])
> training_set[-3] = scale(training_set[-3])
> test_set[-3] = scale(test_set[-3])
> test_set[-3]
```

	Age	EstimatedSalary
2	-0.30419063	-1.51354339
4	-1.05994374	-0.32456026
5	-1.81569686	0.28599864
9	-1.24888202	-1.09579256
12	-1.15441288	-0.48523366
18	0.64050076	-1.32073531
19	0.73496990	-1.25646596
20	0.92390818	-1.22433128
22	0.82943904	-0.58163769
29	-0.87100546	-0.77444577
32	-1.05994374	2.24621408
34	-0.96547460	-0.74231109
35	-1.05994374	0.73588415
38	-0.77653633	-0.58163769
45	-0.96547460	0.54307608
46	-1.43782030	-1.51354339

Step 6: Fitting SVM to the training set

```

> install.packages('e1071')

> library(e1071)
> classifier = svm(formula = Purchased ~ .,
+                  data = training_set,
+                  type = 'C-classification',
+                  kernel = 'linear')
> classifier

Call:
svm(formula = Purchased ~ ., data = training_set, type = "C-classification",
    kernel = "linear")

Parameters:
  SVM-Type:  c-classification
 SVM-Kernel: linear
       cost: 1

Number of Support Vectors: 116

```

Step 7: Predicting the test set result

```

> y_pred = predict(classifier, newdata = test_set[-3])
> y_pred
 2  4  5  9 12 18 19 20 22 29 32 34 35 38 45 46 48 52 66
0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
69 74 75 82 84 85 86 87 89 103 104 107 108 109 117 124 126 127 131
0  0  0  0  0  0  0  0  0  0  0  1  0  0  0  0  0  0  0
134 139 148 154 156 159 162 163 170 175 176 193 199 200 208 213 224 226 228
0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  1  1  1  0  1
229 230 234 236 237 239 241 255 264 265 266 273 274 281 286 292 299 302 305
0  1  1  1  0  1  1  1  0  1  1  1  1  1  1  0  1  1  1  0
307 310 316 324 326 332 339 341 343 347 353 363 364 367 368 369 372 373 380
1  0  0  0  0  1  0  1  0  1  1  0  1  1  1  0  1  0  1
383 389 392 395 400
1  0  0  0  0
Levels: 0 1

```


```

> cm = table(test_set[, 3], y_pred)
> cm
  y_pred
    0  1
0  57  7
1  13 23

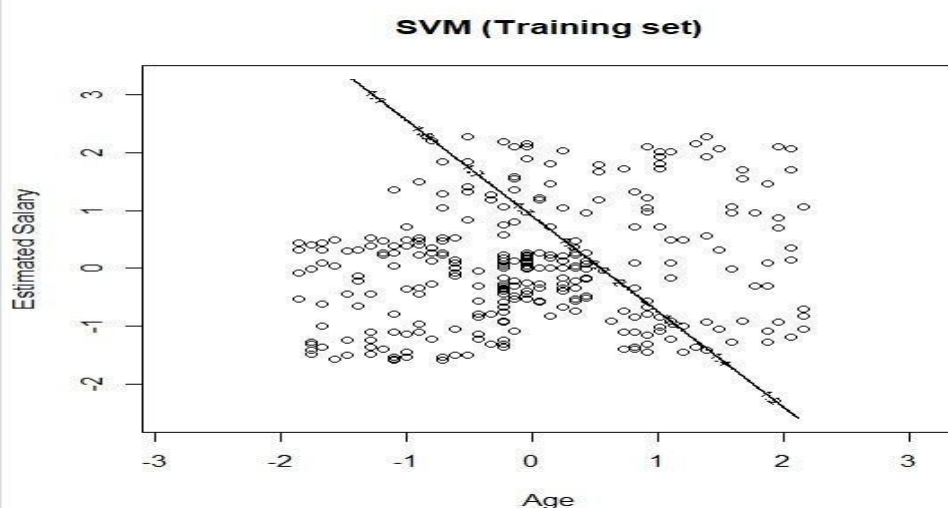
```

Step 8: Visualizing the Training set results

```
> set = training_set
> x1 = seq(min(set[, 1]) - 1, max(set[, 1]) + 1, by = 0.01)
> x2 = seq(min(set[, 2]) - 1, max(set[, 2]) + 1, by = 0.01)
```



```
> grid_set = expand.grid(x1, x2)
> colnames(grid_set) = c('Age', 'EstimatedSalary')
> y_grid = predict(classifier, newdata = grid_set)
> plot(set[, -3],
+       main = 'SVM (Training set)',
+       xlab = 'Age', ylab = 'Estimated salary',
+       xlim = range(x1), ylim = range(x2))
```



```
> contour(x1, x2, matrix(as.numeric(y_grid), length(x1), length(x2)), add = TRUE)  
> points(grid_set, pch = '.', col = ifelse(y_grid == 1, 'coral1', 'aquamarine'))  
> points(set, pch = 21, bg = ifelse(set[, 3] == 1, 'green4', 'red3'))
```



PRACTICAL NO : 4

Aim: Implement an application that stores big data in Hbase / MongoDB and manipulate it using R / Python

ABHISHEK'S ORG - 2024-07-27 > PROJECTS

Create a Project

Name Your Project

Add Members


Name Your Project

Project names have to be unique within the organization (and other restrictions).

Akash|

Add Tags (Optional)

Use tags to efficiently label and categorize your projects. A project can have a maximum of 50 tags. You can modify tags for the project later. [Learn more](#)

Key	Value	Actions
<input type="text" value="Select a key or enter your own"/>	: <input type="text" value="Select a value or enter your own"/>	
+ Add tag		
0 TAGS		

Step 1 : Sign up and create a cluster.

CLUSTERS > CREATE A SHARED CLUSTER

Create a Shared Cluster

Serverless

Dedicated

Shared

For learning and exploring MongoDB in a sandbox environment. Basic configuration controls. No credit card required to start. Upgrade to dedicated clusters for full functionality. Explore with sample datasets. Limit of one free cluster per project.

Cloud Provider & Region AWS, Mumbai (ap-south-1)

aws

Google Cloud

Azure

★ Recommended region ⓘ 🏢 Dedicated tier region ⓘ 🌿 Low carbon emissions region ⓘ

NORTH AMERICA

EUROPE

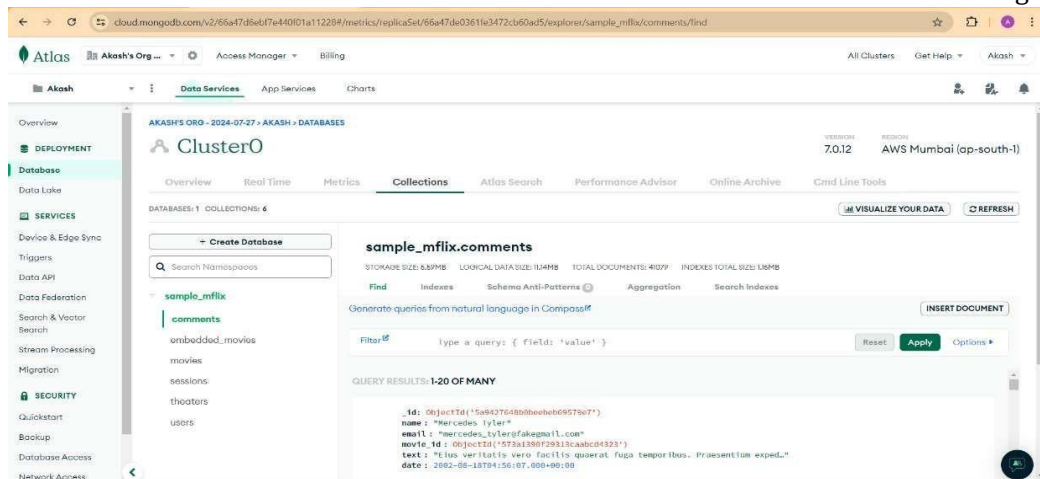
AUSTRALIA

🇺🇸 Oregon (us-west-2) ★ 🌿
🇫🇷 Paris (eu-west-3) ★ 🌿
🇦🇺 Sydney (ap-southeast-2) ★

This is the home page of mongoDB Atlas.

The screenshot shows the MongoDB Atlas interface. The top navigation bar includes 'Atlas', 'Akash's Org', 'Access Manager', and 'Billing'. The main header shows 'Cluster0' with version '7.0.12' and region 'AWS Mumbai (ap-south-1)'. The left sidebar lists various services like 'Database', 'Data Lake', 'Device & Edge Sync', etc. The main content area is titled 'akash_db.akash1' and shows 'STORAGE SIZE: 4KB', 'LOGICAL DATA SIZE: 0B', 'TOTAL DOCUMENTS: 0', and 'INDEXES TOTAL SIZE: 4KB'. The 'Collections' tab is active, displaying a search bar and a 'Filter' section with a query input field. The 'QUERY RESULTS' section shows '0' results.

Step 2 : Click on collections to create and view existing databases.



Step 3 : Click on 'Add My Own Data' to create a database.

The screenshot shows the 'Create Database' dialog box. It has a title 'Create Database' and a close button. Below the title, there are three input fields: 'Database name' with a placeholder 'Enter database name', 'Collection name' with a placeholder 'Enter collection name', and 'Additional Preferences' with a dropdown menu set to 'Select'. At the bottom, there are two buttons: 'Cancel' and 'Create'.

Step 4 : Click on insert document to add records.

The screenshot shows the MongoDB Atlas web interface. The top navigation bar includes the Atlas logo, the organization name 'Akash's Org', and links to 'Access Manager' and 'Billing'. The main header shows 'All Clusters', 'Get Help', and the user 'Akash'. The left sidebar contains navigation options: Overview, Deployment, Database, Data Lake, Services, and Security. The 'Database' section is expanded, showing 'akash_db' and 'akash1'. The main panel displays the 'akash_db.akash1' collection details, including storage size (4KB), logical data size (0B), total documents (0), and index size (4KB). A search bar and a query editor are visible, along with an 'INSERT DOCUMENT' button.

Since MongoDB is a No-SQL database, so you can add 'n' number of columns for any row/record.

Insert Document

To collection akash1

	VIEW	
1	_id : 66a47ef3f1166a4cdb07a0a3	ObjectId
2	name : "AKash"	String
3	id : "7"	String
	city : "mumbai"	String

Cancel

Insert

Perform updating data

Generate queries from natural language in Compass [↗](#) INSERT DOCUMENT

Filter [↗](#) Type a query: { field: 'value' } Reset Apply Options ▼

Project { field: 0 }

Sort { field: -1 } or [['field', -1]]

Collation { locale: 'simple' }

QUERY RESULTS: 1-2 OF 2

1	_id: ObjectId('66a47ef3f1166a4cdb07a0a3')	ObjectId
2	name: 'Abhishek'	String
3	id: '7'	String
4	city: 'mumbai'	String

Document modified. CANCEL UPDATE

Performing deleting data

Generate queries from natural language in Compass [↗](#) INSERT DOCUMENT

Filter [↗](#) Type a query: { field: 'value' } Reset Apply Options ▼

Project { field: 0 }

Sort { field: -1 } or [['field', -1]]

Collation { locale: 'simple' }

QUERY RESULTS: 1-2 OF 2

▶	_id: ObjectId('66a47ef3f1166a4cdb07a0a3')	✎ 📄 🗑️
	name: "Abhishek"	
	id: "7"	
	city: "mumbai"	

_id: ObjectId('66a47fe5f1166a4cdb07a0a4')

Performing Insert data

Charts 👤 👥

STORAGE SIZE: 4KB LOGICAL DATA SIZE: 0B TOTAL DOCUMENTS: 0 INDEXES TOTAL SIZE: 4KB

[Find](#) [Indexes](#) [Schema Anti-Patterns](#) [Aggregation](#) [Search Indexes](#)

Generate queries from natural language in Compass [↗](#) INSERT DOCUMENT

Filter [↗](#) Type a query: { field: 'value' } Reset Apply Options ▼

Project { field: 0 }

Sort { field: -1 } or [['field', -1]]

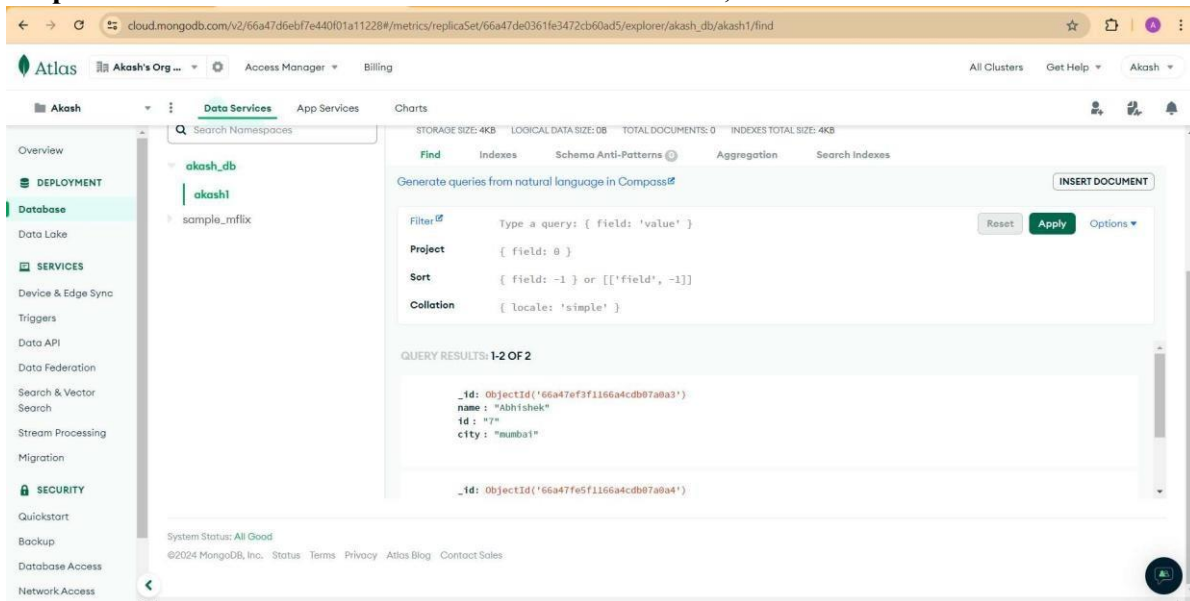
Collation { locale: 'simple' }

QUERY RESULTS: 1-2 OF 2

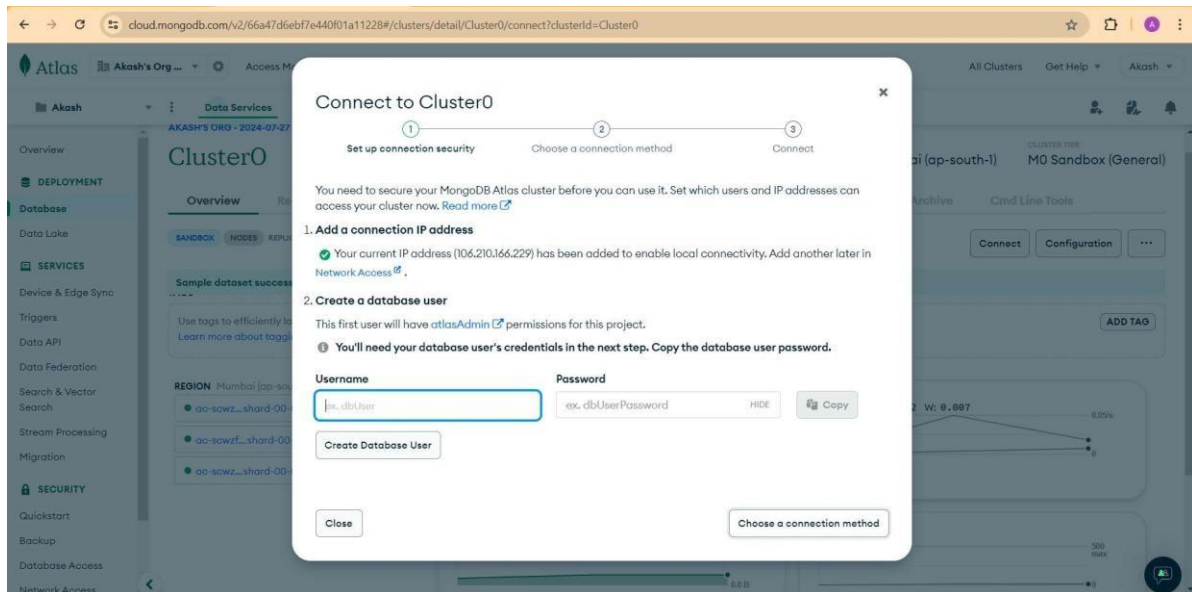
_id: ObjectId('66a47ef3f1166a4cdb07a0a3')
name: "Abhishek"
id: "7"
city: "mumbai"

_id: ObjectId('66a47fe5f1166a4cdb07a0a4')

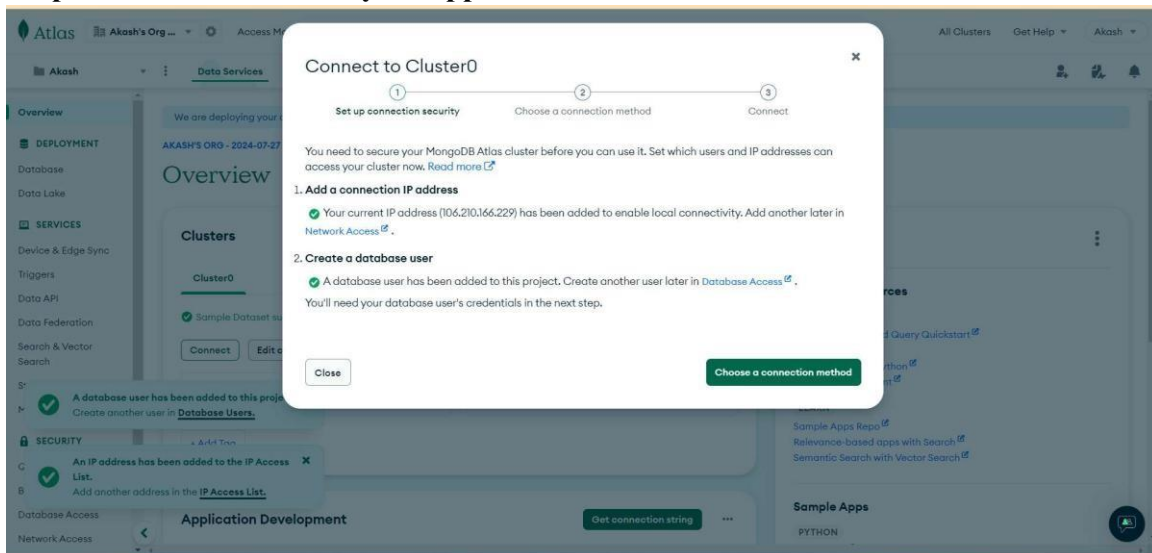
Step 5 : To start with the connection click on Overview, and then click on Connect.



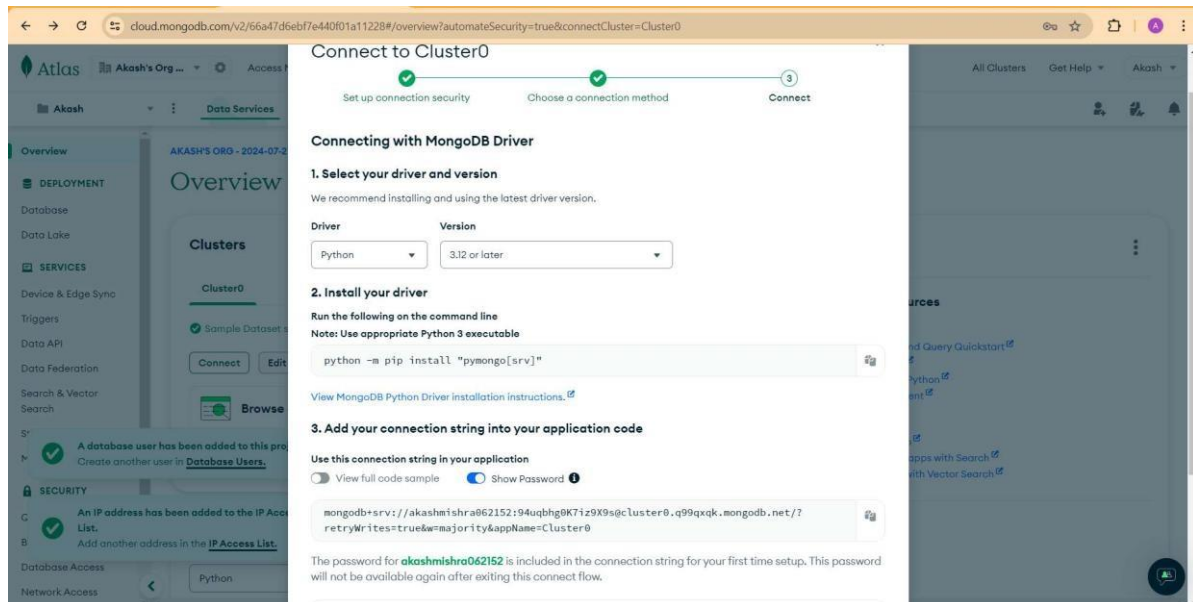
Step 6 : Select on add your current IP and create a MongoDB user.



Step 7 : Click on ‘Connect your application’.



Step 8 : Select the driver as ‘Python’ and version as ‘3.6 or later’. (Select the version as 3.6 or later only if your Python’s version is 3.6 or later.)



Step 9 : Write the code given below in a Python file.

bigdata.py - C:/Users/akash/OneDrive/Desktop/bigdata.py (3.12.4)

File Edit Format Run Options Window Help

```
import pymongo
from pymongo import MongoClient
client=MongoClient("mongodb+srv://akashmishra062152:94uqbhg0K7iz9X9s@cluster0.q99qxqk.mongodb.net/?retryWrites=true&majority=&appName=Cluster0")
db=client.get_database("akash_db")
records=db.akash1
db=client.test
print(records.count_documents({}))
print(list(records.find()))
```

```
>>> ===== RESTART: C:/Users/akash/OneDrive/Desktop/bigdata.py =====
2
[{'_id': ObjectId('66a47fe5f1166a4cdb07a0a4'), 'name': 'supriya ma'am', 'id': '8', 'city': 'mumbai'}, {'_id': ObjectId('66a4805af1166a4cdb07a0a5'), 'name': 'Rahul', 'id': '9', 'city': 'mumbai'}]
>>>
```

Output :

PRACTICAL NO : 5

Aim: write program in R of Naive baye's theorem

Loading data

```
> data(iris)
> str(iris)
'data.frame': 150 obs. of 5 variables:
 $ Sepal.Length: num  5.1 4.9 4.7 4.6 5 5.4 4.6 5 4.4 4.9 ...
 $ Sepal.Width : num  3.5 3 3.2 3.1 3.6 3.9 3.4 3.4 2.9 3.1 ...
 $ Petal.Length: num  1.4 1.4 1.3 1.5 1.4 1.7 1.4 1.5 1.4 1.5 ...
 $ Petal.Width : num  0.2 0.2 0.2 0.2 0.2 0.4 0.3 0.2 0.2 0.1 ...
 $ Species      : Factor w/ 3 levels "setosa","versicolor",...: 1 1 1 1
```

Installing Packages

```
> install.packages("e1071")
> install.packages("caTools")
> install.packages("caret")
```

Loading package

```
> library(e1071)
> library(caTools)
> library(caret)
Loading required package: lattice
Loading required package: ggplot2
```

Splitting data into train and test data

```
> split <- sample.split(iris, splitRatio = 0.7)
> train_cl <- subset(iris, split == "TRUE")
> test_cl <- subset(iris, split == "FALSE")
>
> train_scale <- scale(train_cl[, 1:4])
> test_scale <- scale(test_cl[, 1:4])
>
> set.seed(120) # Setting seed
> classifier_cl <- naiveBayes(Species ~ ., data = train_cl)
> classifier_cl
```

Naive Bayes Classifier for Discrete Predictors

call:

naiveBayes.default(x = x, y = y, laplace = laplace)

A-priori probabilities:

```
Y
  setosa versicolor virginica
0.3333333 0.3333333 0.3333333
```

Conditional probabilities:

```
      Sepal.Length
Y      [,1]      [,2]
setosa  5.046667 0.3848272
versicolor 5.963333 0.5268536
virginica  6.553333 0.6693967
```

```
      Sepal.Width
Y      [,1]      [,2]
setosa  3.413333 0.4256705
versicolor 2.823333 0.3470897
virginica  2.956667 0.3136914
```

```
      Petal.Length
Y      [,1]      [,2]
setosa  1.466667 0.1561019
versicolor 4.320000 0.4759020
virginica  5.496667 0.5738457
```

```
      Petal.Width
Y      [,1]      [,2]
setosa  0.2766667 0.1135124
versicolor 1.3533333 0.1960530
virginica  2.0433333 0.2568823
```

Predicting on test data'

```
> y_pred <- predict(classifier_cl, newdata = test_cl)
> cm <- table(test_cl$species, y_pred)
> cm
```

	y_pred		
	setosa	versicolor	virginica
setosa	20	0	0
versicolor	0	19	1
virginica	0	2	18

```
>
```

Model Evauation

```
> confusionMatrix(cm)
Confusion Matrix and Statistics
```

	y_pred		
	setosa	versicolor	virginica
setosa	20	0	0
versicolor	0	19	1
virginica	0	2	18

```
overall statistics
```

```
Accuracy : 0.95
95% CI : (0.8608, 0.9896)
No Information Rate : 0.35
P-Value [Acc > NIR] : < 2.2e-16
```

```
Kappa : 0.925
```

```
Mcnemar's Test P-Value : NA
```

```
Statistics by Class:
```

	Class: setosa	Class: versicolor	Class: virginica
sensitivity	1.0000	0.9048	0.9474
specificity	1.0000	0.9744	0.9512
Pos Pred Value	1.0000	0.9500	0.9000
Neg Pred Value	1.0000	0.9500	0.9750
Prevalence	0.3333	0.3500	0.3167
Detection Rate	0.3333	0.3167	0.3000
Detection Prevalence	0.3333	0.3333	0.3333
Balanced Accuracy	1.0000	0.9396	0.9493

PRACTICAL NO : 6

Aim: WAP showing implementation of Regression model.

Regression is a method to mathematically formulate relationship between variables that in due course can be used to estimate, interpolate and extrapolate. Suppose we want to estimate the weight of individuals, which is influenced by height, diet, workout, etc.

Here, *Weight* is the **predicted** variable


Lets implementation of Regression Model some Example:

```
> IQ <- rnorm(40, 30, 2)
> IQ <- sort(IQ)


> result <- c(0, 0, 0, 1, 0, 0, 0, 0, 0, 1,
+ 1, 0, 0, 0, 1, 1, 0, 0, 1, 0,
+ 0, 0, 1, 0, 0, 0, 1, 1, 0, 1, 1,
+ 1, 1, 1, 0, 1, 1, 1, 1, 0, 1)

> df <- as.data.frame(cbind(IQ, result))
> print(df)
      IQ result
1 25.58824     0
2 26.43200     0
3 27.37083     0
4 27.37898     1
5 27.56671     0
6 28.08275     0
7 28.35637     0
8 28.41538     0

> png(file="LogisticRegressionGFG.png")
> plot(IQ, result, xlab = "IQ Level",
+ ylab = "Probability of Passing")
> g = glm(result~IQ, family=binomial, df)
```



```
> curve(predict(g, data.frame(IQ=x), type="resp"), add=TRUE)  
> points(IQ, fitted(g), pch=30)
```




```

> summary(g)

Call:
glm(formula = result ~ IQ, family = binomial, data = df)

Deviance Residuals:
    Min       1Q   Median       3Q      Max
-1.9877 -0.9804 -0.4502  0.9731  1.8898

Coefficients:
              Estimate Std. Error z value Pr(>|z|)
(Intercept) -14.4934     5.8835  -2.463  0.0138 *
IQ            0.4708     0.1922   2.450  0.0143 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

    Null deviance: 55.352  on 39  degrees of freedom
Residual deviance: 47.090  on 38  degrees of freedom
AIC: 51.09

Number of Fisher Scoring iterations: 4

> dev.off()
null device
      1

```


PRACTICAL NO : 7

Aim: WAP showing clustering.

Step 1: Apply kmeans to *newiris*, and store the clustering result in *kc*. The clusternumber is set to 3.

```
> newiris <- iris
> newiris$species <- NULL
> (kc <- kmeans(newiris, 3))
K-means clustering with 3 clusters of sizes 38, 62, 50

Cluster means:
  Sepal.Length Sepal.Width Petal.Length Petal.Width
1    6.850000    3.073684    5.742105    2.071053
2    5.901613    2.748387    4.393548    1.433871
3    5.006000    3.428000    1.462000    0.246000

Clustering vector:
 [1] 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3
[35] 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3 2 2 1 2 2 2 2 2 2 2 2 2 2 2 2 2
[69] 2 2 2 2 2 2 2 2 2 1 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 2
[103] 1 1 1 1 2 1 1 1 1 1 1 2 2 1 1 1 1 2 1 2 1 2 1 1 2 2 1 1 1 1 1
[137] 1 1 2 1 1 1 2 1 1 1 2 1 1 2

within cluster sum of squares by cluster:
[1] 23.87947 39.82097 15.15100
(between_ss / total_ss = 88.4 %)
```

Available components:

```
[1] "cluster"      "centers"      "totss"        "withinss"
[5] "tot.withinss" "betweenss"    "size"         "iter"
[9] "ifault"
```

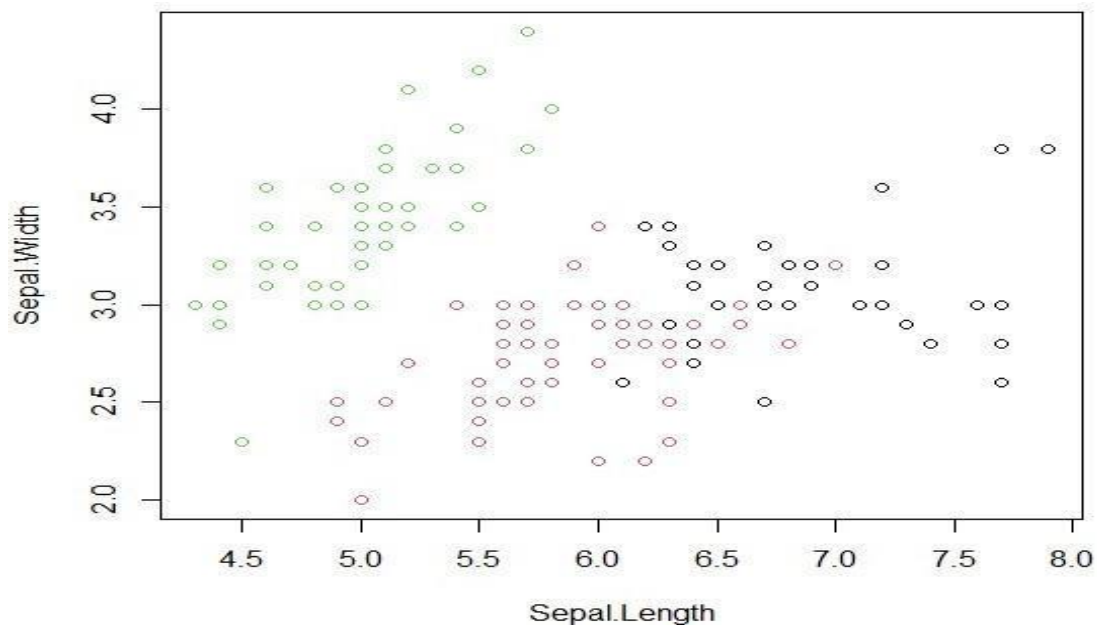
Step 2: Compare the Species label with the clustering result

```
> table(iris$species, kc$cluster)

      1  2  3
setosa  0  0 50
versicolor  2 48  0
virginica 36 14  0
```


Step 3 : Plot the clusters and their centres. Note that there are four dimensions in the data and that only the first two dimensions are used to draw the plot below.

```
> plot(newiris[c("Sepal.Length", "Sepal.Width")], col=kc$cluster)
```

Step 4: Some black points close to the green centre (asterisk) are actually closer to the black centre in the four dimensional space.

```
> points(kc$centers[,c("Sepal.Length", "Sepal.Width")], col=1:3, pch=8, cex=2)
```



PRACTICAL NO : 8

Multiple regression

Aim: Apply Multiple regressions, if data have a continuous independent variable .

```
> install.packages("tidyverse")
Installing package into 'C:/Users/praja/AppData/Local/R/win-library/4.4'
(as 'lib' is unspecified)
--- Please select a CRAN mirror for use in this session ---
also installing the dependencies 'fastmap', 'colorspace', 'sys', 'bit', 'ps', 'base64enc',

There are binary versions available but the source versions are later:
      binary source needs_compilation
colorspace 2.1-0 2.1-1              TRUE
yaml       2.3.9 2.3.10             TRUE

> library(tidyverse)
> data("marketing", package = "datarium")
Error in find.package(package, lib.loc, verbose = verbose) :
  there is no package called 'datarium'
> install.packages("datarium")
Installing package into 'C:/Users/praja/AppData/Local/R/win-library/4.4'
(as 'lib' is unspecified)
trying URL 'https://cran.icts.res.in/bin/windows/contrib/4.4/datarium_0.1.0.z
Content type 'application/zip' length 48431 bytes (47 KB)
downloaded 47 KB

package 'datarium' successfully unpacked and MD5 sums checked

The downloaded binary packages are in
      C:\Users\praja\AppData\Local\Temp\RtmpkrAieH\downloaded_packages
> data("marketing", package = "datarium")
> head(marketing, 4)
  youtube facebook newspaper sales
1  276.12    45.36    83.04 26.52
2   53.40    47.16    54.12 12.48
3   20.64    55.08    83.16 11.16
4  181.80    49.56    70.20 22.20
> model <- lm(sales ~ youtube + facebook + newspaper, data = marketing)
> summary(model)
```

```

> summary(model)$coefficient
              Estimate Std. Error    t value    Pr(>|t|)
(Intercept)  3.526667243 0.374289884   9.4222884 1.267295e-17
youtube      0.045764645 0.001394897  32.8086244 1.509960e-81
facebook     0.188530017 0.008611234  21.8934961 1.505339e-54
newspaper    -0.001037493 0.005871010  -0.1767146 8.599151e-01
> model <- lm(sales ~ youtube + facebook, data = marketing)
> summary(model)


Call:
lm(formula = sales ~ youtube + facebook, data = marketing)

Residuals:
    Min       1Q   Median       3Q      Max
-10.5572  -1.0502   0.2906   1.4049   3.3994

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  3.50532     0.35339   9.919  <2e-16 ***
youtube      0.04575     0.00139  32.909  <2e-16 ***
facebook     0.18799     0.00804  23.382  <2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.018 on 197 degrees of freedom
Multiple R-squared:  0.8972,    Adjusted R-squared:  0.8962
F-statistic: 859.6 on 2 and 197 DF,  p-value: < 2.2e-16

```



```

> confint(model)
              2.5 %      97.5 %
(Intercept) 2.80841159 4.20222820
youtube     0.04301292 0.04849671
facebook    0.17213877 0.20384969

```

